08 — Visual inconsistency review & AI-generation assessment

What it detects: whether a page *looks* edited or synthetic to a vision model — inconsistent fonts, misaligned text, patchy colour/contrast, uneven sharpness or lighting — and whether it appears AI-generated. It is a **secondary, judgment-based** signal that can catch what pixel forensics miss, or corroborate them.

Code: `services/visual_inconsistency_service.py` (`run_visual_inconsistency_review`), `services/llm_service.py` (`analyze_page_visual_consistency`, `PageVisualAnalysis`), run by `tasks/visual_inconsistency_task.py`; stores `check_type = visual_inconsistency_review`.

Algorithm

1. Render every page of the PDF (`render_pdf_pages`, 200 DPI) and encode as PNG, longest side capped at **1536 px** (`MAX_IMAGE_DIMENSION`). PNG, not JPEG, so glyph edges stay crisp and compression artifacts are not introduced.
2. Send each page to the vision model **twice, independently** (`RUNS_PER_PAGE = 2`), at `temperature = 0.4` — deliberately non-zero, unlike every other call in `LLMService`, so the runs are genuinely independent samples. Each run answers six structured questions: font consistency, text alignment, colour/contrast, resolution/sharpness, shadow/lighting — each a `VisualCategoryFinding {consistent, description, confidence, bounding_box}` — plus the `AIGenerationAssessment {likely_ai_generated, description, confidence}`.
3. **Corroboration rule** (the reliability mechanism): a category is reported as a real finding only if **both runs** flagged it, or **either** run flagged it with `high` confidence. A single low/medium one-off is kept but labelled "treated as noise" at severity `info`.
4. **Severity:** for a confirmed finding, `high` if both runs agreed *and* at least one was high-confidence, else `medium`; everything not confirmed (both runs consistent, or an uncorroborated one-off) is `info`. The check's `result` is `flag` if any finding is `medium` or `high`, else `pass`.
5. The task appends `metadata_forensics_correlation` (info) cross-referencing the metadata result, so a reviewer sees editing-tool evidence next to visual evidence. It is **not scored** (avoids double counting), as is `experimental_signal_notice` (an info notice always added to the result).

Findings

`finding`	Meaning
`visual_font_consistency`	Fonts vary within the page
`visual_text_alignment`	Text or table alignment is inconsistent
`visual_color_contrast_consistency`	Colour/contrast differs across regions
`visual_resolution_sharpness_consistency`	Sharpness differs across regions
`visual_shadow_lighting_consistency`	Shadows/lighting inconsistent
`ai_generation_assessment`	The page may be AI-generated/synthetic
`experimental_signal_notice`, `metadata_forensics_correlation`	Info only, never scored

Output shape (real, truncated)

```
{
  "result": "flag",                          // flag if any corroborated finding, else pass
  "details": [
    {"finding": "experimental_signal_notice", "severity": "info",
     "description": "This check uses a vision-capable language model's visual judgment, run twice
independently…"},
    {"finding": "visual_font_consistency", "severity": "high", "page": 1,
     "description": "Page 1 — Font consistency: The font weight and style for the item descriptions … This
is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding.",
     "bounding_box": {"page": 1, "x": 0.05, "y": 0.3, "width": 0.9, "height": 0.15},
     "data": {
       "category": "font_consistency",
       "run_agreement": "2-of-2 runs",       // | "1-of-2 runs (high confidence)" | "consistent in both
runs" | "1-of-2 runs, not corroborated" | "not flagged"
       "run_1": {"consistent": false, "confidence": "high", "description": "…", "bounding_box": {"x": 0.05,
"y": 0.3, "width": 0.9, "height": 0.15}},
       "run_2": {"consistent": false, "confidence": "high", "description": "…", "bounding_box": {…}}}}
  ]
}
```

`data` keeps **both raw runs**, so a reviewer or developer can see exactly what was said and why a finding did or did not survive. The `bounding_box` (from the first run that supplied one) is drawn as a **dashed blue "AI-described area (approximate)"** overlay — deliberately styled to look less precise than the red/orange pixel-forensic boxes.

Known limitations

- **It is a language model's opinion.** Runs vary; corroboration reduces but does not remove hallucination. Boxes are approximate (models are poor at localisation).
- **AI-generation detection is unproven** on scanned/business documents. It is weighted lowest of any visual signal (weight 6) and every description says "experimental, probabilistic".
- **Every page costs two vision calls.** A 10-page PDF is 20 calls; there is no page cap on this task (unlike signature detection, capped at 10 pages).
- Legitimate documents made of mixed sources (a template with a pasted logo, mixed fonts) can flag.
- PDF only.

Risk rules fed

Each rule fires only on `medium` / `high` findings (`severity_in`).

`rule_id`	Finding	Weight	Severity
`visual.font_inconsistency`	`visual_font_consistency`	12	medium
`visual.alignment_inconsistency`	`visual_text_alignment`	10	medium
`visual.color_contrast_inconsistency`	`visual_color_contrast_consistency`	12	medium
`visual.sharpness_inconsistency`	`visual_resolution_sharpness_consistency`	10	medium
`visual.lighting_inconsistency`	`visual_shadow_lighting_consistency`	8	low
`ai.generated_content_suspected`	`ai_generation_assessment`	6	low